

Semaphores and FreeRTOS

ENCE464 Embedded Software and Advanced Computing

Course coordinator: Steve Weddell (steve.weddell@canterbury.ac.nz)

Lecturer: Le Yang (le.yang@canterbury.ac.nz)

Department of Electrical and Computer Engineering

Where we're going today

- **Introduction to semaphores**
- Signaling with semaphores
- Mutex and Queues

Introduction to Semaphores

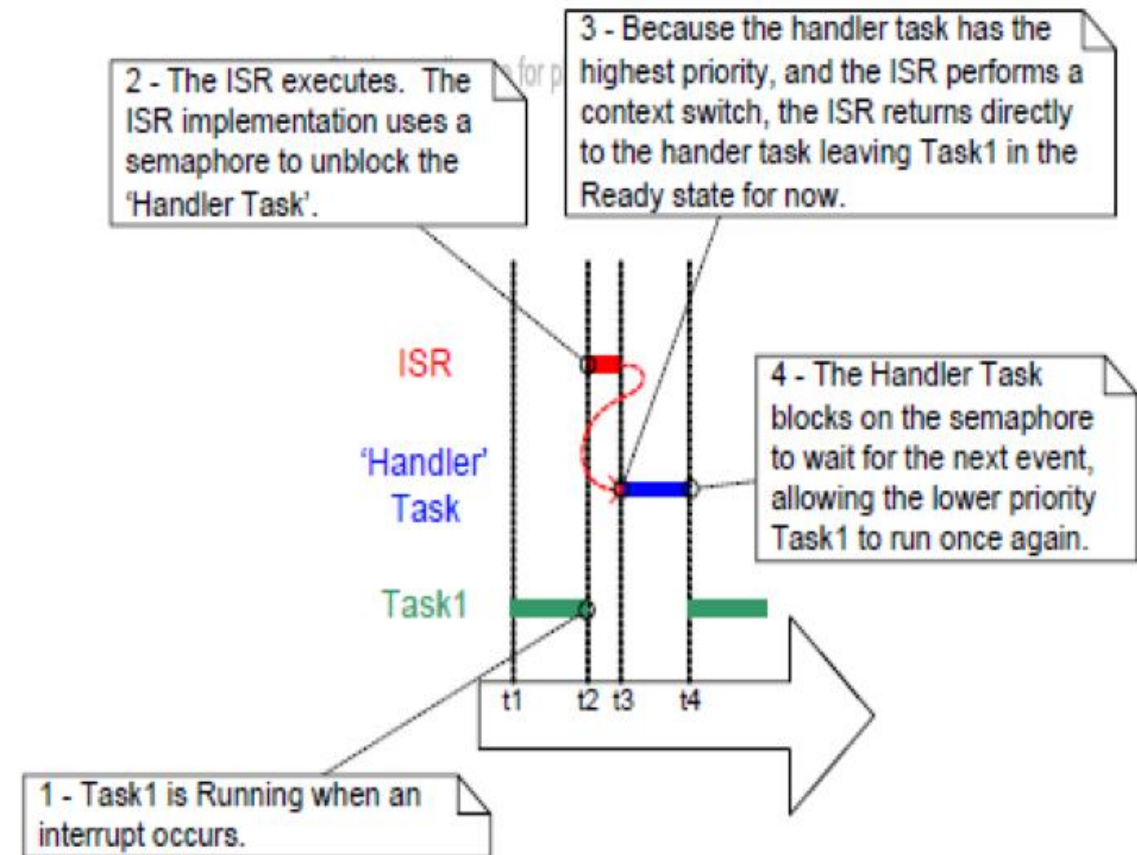
- Semaphore is like an integer with three differences
 - When a semaphore is created, it can be initialized to *any* integer
 - After that only **increment (increase by 1)** and **decrement (decrease by 1)** can be performed via **functions**
 - If the result, after a task decrements a semaphore, is negative, the task blocks itself until another task increments the semaphore
 - The task notifies the scheduler (kernel) that it cannot proceed until an event causes the task to become unblocked
 - When a task increments a semaphore, if there are other tasks waiting, one of the waiting tasks become unblocked

Language of Semaphores

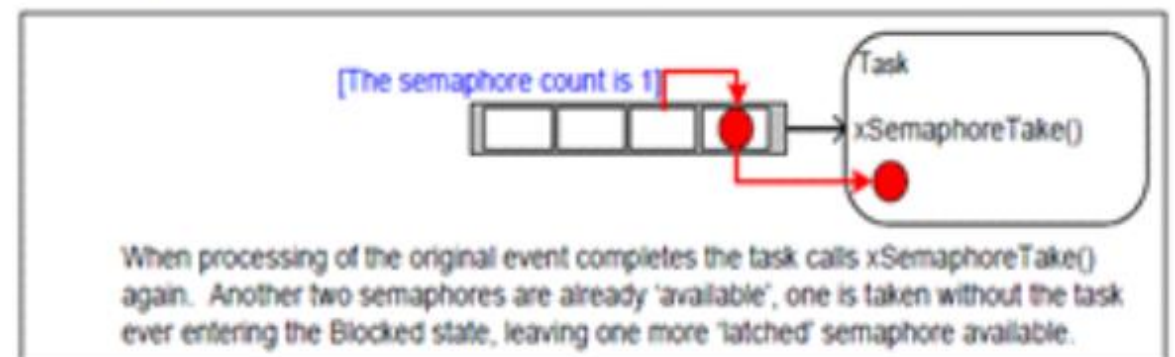
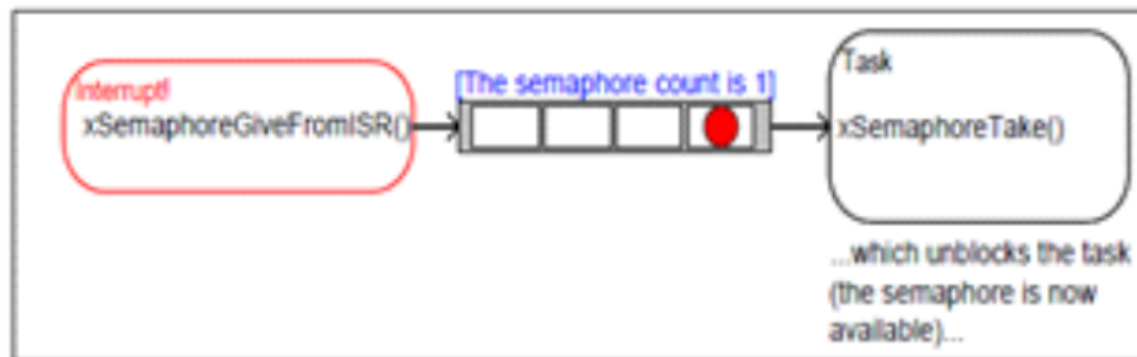
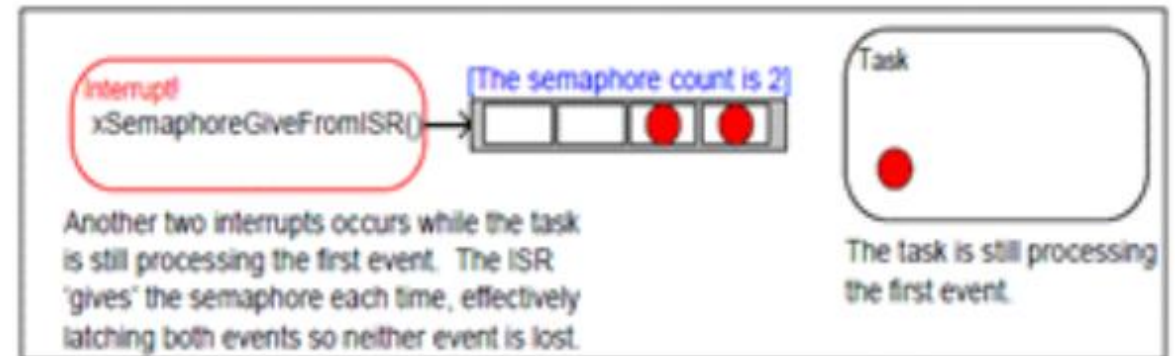
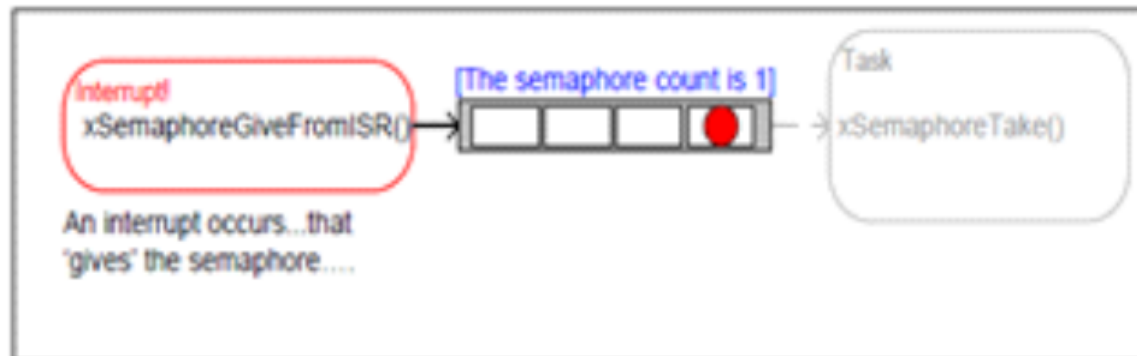
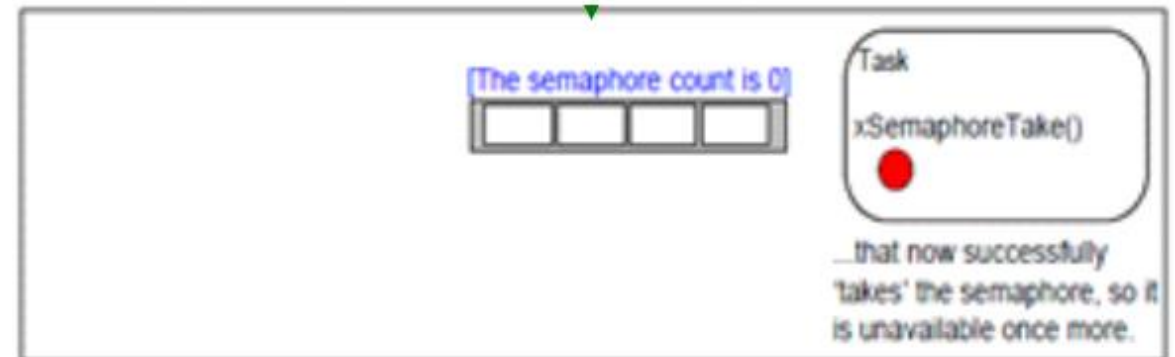
- increment(), decrement()
- signal(), wait()
- FreeRTOS
 - xSemaphoreGive(), xSemaphoreTake()
 - xSemaphoreGiveFromISR()
 - xSemaphoreGiveRecursive(), xSemaphoreTakeRecursive()

Example: Deferred Interrupt Structure

- Use a *binary* semaphore to unblock a task each time a particular interrupt occurs
 - Allows the majority of ISR to be realized within the task
 - Interrupt service 'deferred' to a 'handler' task
- 'Handler' task can be set to have high priority so that it can preempt other tasks
 - ISR can include a context switch to ensure that ISR returns to 'handler' task
 - Entire processing runs contiguously



Example: Counting Semaphore



Why Using Semaphores?

- Semaphores impose deliberate constraints that can help programmers avoid errors
- Solutions using semaphores are often clean and organized, making it easy to check their correctness
- Semaphores can be implemented efficiently on many systems
 - Solutions using semaphores are portable

Where we're going today

- Introduction to semaphores
- **Signaling with semaphores**
- Mutex and Queues

Signaling

- Possibly the simplest use of a semaphore
 - A task sends a signal to another task to indicate that an event has occurred
 - Make it possible a section of codes in one task will run before a section of codes in another task
- Example
 - sem initialized to be 0
 - The order of a1 followed by b1 is guaranteed

Thread A	
1	statement a1
2	sem.give()

Thread B	
1	sem.take()
2	statement b1

Signaling with a Semaphore (1)

- Requirement
 - a1 happens before b2
 - b1 happens before a2

Thread A	
1	statement a1
2	statement a2

Thread B	
1	statement b1
2	statement b2

Signaling with a Semaphore (2)

- Is this correct?

Thread A	
1	statement a1
2	bArrived.take()
3	aArrived.give()
4	statement a2

Thread B	
1	statement b1
2	aArrived.take()
3	bArrived.give()
4	statement b2

- **Deadlock!**

Where we're going today

- Introduction to semaphores
- Signaling with semaphores
- **Mutex and Queues**

Mutexes

- Mutexes are used to achieve mutual exclusion
 - Control concurrent access to shared resources (e.g., a buffer or a critical variable)
 - Guarantee that only one task accesses the shared variable at a time
- A mutex is like a token that passes from one task to another, allowing one task at a time to proceed
- In order for a task to access a shared variable, it has to obtain mutex first and it releases the mutex when it is done

Queues

- A queue in FreeRTOS holds a finite number of fixed-size data items
 - Maximum number of items a queue can hold is called its 'length'
 - First In First Out (FIFO) buffer
- Access by multiple tasks
 - Any number of tasks can write to and read from a queue
 - A queue having multiple writers is common
 - A queue having multiple readers is rare

Case Study (1)

- Flight data recorder (FDR)
 - “Blackbox”
 - Record important flight parameters
 - Control and actuator positions
 - Engine information
 - Time of day
 - Minimum 88 parameters under current U.S. federal regulations
 - Each parameter is recorded a few times per second
 - But bursts of data at a much higher frequency are possible
 - Most FDR record approximately 17-25 hours of data in a continuous loop
 - Annual FDR check (readout) is performed annually



Case Study (2)

- Tasks suggested
 - Maintain a real-time clock (F)
 - Monitor signal changes (F)
 - Go into low-power mode on loss of external power (F)
 - Detect unauthorized access attempt (F)
- Record samples into SRAM buffer
- Compress and store data in flash memory
- Maintain status display
- Monitor and report battery status
- Dump all data to flash memory
- Allow download on recovery
- ...